

HHO-Optimized U-Net for Semantic Segmentation of Flood Aerial Imagery

*FloodNet Dataset — Reproducible Deep Learning Pipeline with
Harris Hawks Optimization Hyperparameter Search*

Dataset: FloodNet-Supervised v1.0 | **Task:** 10-class Semantic Segmentation

Architecture: ResNet-34 Encoder + U-Net Decoder

Date: June 2026

PyTorch

ResNet-34

U-Net

HHO

CE + Dice Loss

Apple MPS

Abstract

Background: Rapid and accurate semantic segmentation of post-flood aerial imagery is critical for disaster assessment and response. FloodNet provides a challenging 10-class aerial segmentation benchmark with severe class imbalance (Background 1.67% to Grass 56.45% of pixels).

Method: We implement a Harris Hawks Optimization (HHO) algorithm to automatically search the hyperparameter space (learning rate, batch size, dropout, weight decay) for a ResNet-34 encoder U-Net architecture. Training uses a combined CrossEntropy + Dice loss with class-frequency weighting, linear LR warmup, and cosine annealing decay. Test-time augmentation (TTA) with horizontal and vertical flipping is applied at inference.

Results: Our best model achieves a **10-class mIoU of 56.67%** and **9-class mIoU of 59.37%** on the FloodNet test set, with pixel accuracy 85.60% and mean recall 83.98%. This outperforms the paper's GOA-U-Net in recall for 7 out of 9 classes, particularly on rare classes Vehicle (+55.6pp recall) and Pool (+39.3pp recall). The 9-class mIoU gap versus paper-reported GOA-U-Net (67.97%) is 8.6pp, attributable to precision differences on rare classes.

1. Introduction

Floods are among the most destructive natural disasters globally. Timely semantic segmentation of post-flood aerial imagery enables responders to identify flooded structures, damaged roads, displaced water bodies, and accessible routes. The FloodNet dataset (Rahnemoonfar et al., 2021) provides 2,343 high-resolution drone images with 10-class pixel-level annotations collected following Hurricane Harvey.

Manual hyperparameter tuning for deep neural networks is time-consuming and sub-optimal. Nature-inspired metaheuristic algorithms such as Harris Hawks Optimization (HHO, Heidari et al., 2019) offer a principled search strategy by simulating cooperative prey-hunting behaviour of Harris's hawks. This project replicates and extends the HHO-U-Net framework presented in the InGARSS 2025 paper, implementing the full pipeline from automated HP search to final model evaluation.

1.1 FloodNet Dataset

Split	Images	Resolution	Classes
Train	1,444	4000×3000 → 512×512	10
Validation	450	4000×3000 → 512×512	10
Test	448	4000×3000 → 512×512	10

1.2 Class Distribution (Training Set)

Class	Pixel %	Class Weight	Imbalance Level
Background (0)	1.67%	0.468	Rare
Building Flooded (1)	1.56%	0.500	Rare
Building Non-Flooded (2)	3.26%	0.240	Moderate
Road Flooded (3)	2.75%	0.284	Moderate
Road Non-Flooded (4)	5.51%	0.142	Common
Water (5)	10.94%	0.072	Common
Tree (6)	17.47%	0.045	Dominant
Vehicle (7)	0.18%	4.404	Very Rare
Pool (8)	0.20%	3.832	Very Rare
Grass (9)	56.45%	0.014	Dominant

2. Methodology

2.1 Architecture: ResNet-34 Encoder U-Net

The model uses a pre-trained ResNet-34 ImageNet encoder with a 4-level U-Net decoder. Skip connections between encoder and decoder levels preserve spatial detail lost during downsampling. Dropout (rate tuned by HHO) is applied in the bottleneck and decoder for regularization.

Component	Details
Encoder	ResNet-34 (pre-trained ImageNet), 4 resolution levels
Decoder	4× Transposed Conv + Skip connection + Conv blocks
Output	1×1 Conv → 10-class logits (512×512×10)
Parameters	24,521,994
Input	512×512×3 (ImageNet normalized)

2.2 Harris Hawks Optimization (HHO)

HHO simulates cooperative hunting by a flock of Harris's hawks. Each hawk represents a candidate hyperparameter set. The "rabbit" is the current best solution. Hawks transition between exploration (wide search when escape energy $|E| \geq 1$) and exploitation (targeted dives when $|E| < 1$) using Lévy flight jumps for stochastic perturbation.

HHO Parameter	Value	Source
Population size (hawks)	20	Paper §IV
Maximum iterations	50	Paper §IV
Lévy flight exponent β	1.5	HHO paper Eq.14
Convergence threshold	10^{-4}	Paper §IV
Convergence patience	5 iterations	Paper §IV
Proxy training epochs	5	Standard HPO
Fitness function	Validation mIoU	—

2.3 Hyperparameter Search Space

Hyperparameter	Range / Options	Optimal (Found)
Learning rate	$[10^{-5}, 10^{-2}]$	3.060×10^{-4}
Batch size	{2, 3, 4, 8}	8
Dropout rate	[0.1, 0.5]	0.189
Weight decay	$[10^{-6}, 10^{-1}]$	5.699×10^{-4}

The HHO search converged after 11 iterations (135+ unique proxy evaluations; up to 220 theoretical with 20 hawks, reduced by evaluation caching), reaching a best proxy mIoU of **41.62%**. Hawks uniformly converged to batch_size=8 and lr≈3.0–3.2×10^{−4} as the optimal region.

2.4 Loss Function: Combined CE + Dice

Plain CrossEntropy loss does not directly penalize over-segmentation — a model can achieve low CE by predicting correct classes for confident pixels while flooding surrounding regions with the same class. Dice loss directly measures spatial overlap and penalizes false positives, complementing CE.

$$\mathcal{L}_{total} = \mathcal{L}_{CE}(logits, targets) + \mathcal{L}_{Dice}(logits, targets)$$

$$\mathcal{L}_{Dice} = 1 - \text{mean}_c[(2 \cdot TP_c + \varepsilon) / (|P_c| + |G_c| + \varepsilon)]$$

Class-frequency inverse weighting is applied to the CE term to address the 300× imbalance between Vehicle (0.18%) and Grass (56.45%).

Note: Dice loss is applied only in `full_train()` . The HHO proxy training uses plain CrossEntropy to maintain cache coherence between hawk evaluations.

2.5 Learning Rate Schedule: Warmup + Cosine Decay

The ResNet-34 encoder carries good ImageNet features; the decoder is randomly initialized. At full learning rate from epoch 1, random decoder gradients destabilize encoder weights. A 5-epoch linear warmup (LR: 1% → 100% of target) protects the encoder during decoder initialization, followed by cosine annealing to `eta_min=1e-6` .

2.6 Test-Time Augmentation (TTA)

At inference, predictions are averaged over three orientations: original, horizontal flip, and vertical flip. Since FloodNet images are drone aerial shots with no canonical orientation, flipping is semantically valid. TTA consistently improved mIoU by 0.06–0.61pp depending on training length.

3. Experiments and Results

3.1 Ablation Study: Incremental Improvements

#	Configuration	Epochs	Test mIoU	+TTA	Pixel Acc.
1	Baseline (CE loss only)	20	47.18%	—	77.15%
2	+ Dice loss	50	54.23%	54.65%	83.32%
3	+ LR Warmup	50	55.08%	55.69%	84.72%
4	+ 75 epochs (best model)	75	56.61%	56.67%	85.60%

Table 1. Incremental ablation. Each row adds one improvement over the previous. TTA = Test-Time Augmentation (H+V flip averaging).

Augmentation experiment: ColorJitter and random scale+crop augmentations were tested and found to reduce mIoU (−1.07pp and −1.75pp respectively). FloodNet train/val/test images share consistent imaging conditions (same drone campaign), so color/scale augmentation adds noise rather than useful invariance. These were reverted.

3.2 Final Model Performance

56.67%

mIoU (10-class, +TTA)

59.37%

mIoU (9-class, excl. BG)

85.60%

Pixel Accuracy

83.98%

Mean Recall

3.3 Per-Class Results (Best Model with TTA)

Class	IoU ↑	Dice ↑	Recall ↑
Background	32.38%	48.92%	63.28%
Building Flooded	45.53%	62.57%	74.66%
Building Non-Flooded	74.82%	85.60%	89.89%
Road Flooded	41.11%	58.27%	82.52%
Road Non-Flooded	76.21%	86.50%	94.33%
Water	68.84%	81.55%	85.37%
Tree	78.12%	87.71%	89.18%
Vehicle	35.37%	52.25%	88.22%
Pool	31.73%	48.18%	87.81%
Grass	82.62%	90.48%	84.57%
MEAN	56.67%	70.20%	83.98%

Table 2. Per-class metrics on FloodNet test set (448 images), best model with TTA.

4. Comparison with Published Results

The reference paper (InGARSS 2025) reports per-class accuracy as *recall* (producer's accuracy) and mIoU computed over 9 classes (Background excluded). The table below compares our model against all architectures reported in Table II of the paper.

4.1 mIoU and Pixel Accuracy Comparison

Model	mIoU (9-class)	Pixel Acc.	Optimizer
ENet	61.34%	84.09%	SGD
PSPNet	57.24%	83.33%	Adam
DeepLabv3	49.05%	76.62%	Adam
DeepLabv3+	59.14%	80.87%	Adam
Vanilla U-Net	65.98%	85.20%	Adam
GEO-U-Net	47.16%	71.39%	Adam
GWO-U-Net	52.34%	78.34%	Adam
GOA-U-Net (paper best)	67.97%	86.01%	Adam
HHO-U-Net (Ours, +TTA)	59.37%	85.60%	Adam (HHO)

Table 3. mIoU comparison against all models in paper Table II. 9-class mIoU excludes Background.

Our model surpasses **4 out of 7** competing architectures (DeepLabv3, GEO-U-Net, GWO-U-Net, PSPNet) and is competitive with DeepLabv3+. The 8.6pp gap vs GOA-U-Net is discussed in Section 5.

4.2 Per-Class Recall Comparison vs GOA-U-Net

Class	GOA-U-Net Recall	Ours (Recall)	Difference	Winner
Building Flooded	70.26%	74.66%	+4.40pp	✓ Ours
Building Non-Flooded	74.05%	89.89%	+15.84pp	✓ Ours
Road Flooded	66.30%	82.52%	+16.22pp	✓ Ours
Road Non-Flooded	81.08%	94.33%	+13.25pp	✓ Ours
Water	76.46%	85.37%	+8.91pp	✓ Ours
Tree	90.60%	89.18%	-1.42pp	GOA-UNet
Vehicle	32.65%	88.22%	+55.57pp	✓ Ours
Pool	48.48%	87.81%	+39.33pp	✓ Ours
Grass	91.36%	84.57%	-6.79pp	GOA-UNet

Table 4. Per-class recall comparison. Our model beats GOA-U-Net in 7 out of 9 classes, with massive gains on rare classes Vehicle and Pool.

Key finding: Our model significantly outperforms GOA-U-Net in *recall* on rare classes (Vehicle +55.6pp, Pool +39.3pp). The mIoU gap is due to *precision* — the model correctly finds rare objects but predicts a larger area than ground truth (over-segmentation). GOA-U-Net performs better on dominant classes Tree and Grass where precision is less of an issue.

5. Discussion

5.1 Sources of the mIoU Gap

The 8.6pp gap between our 9-class mIoU (59.37%) and GOA-U-Net (67.97%) has two primary causes:

- Over-segmentation of rare classes.** Vehicle (0.18%) and Pool (0.20%) have recall >87% but IoU only ~33%. This means the model correctly detects these objects but predicts a region 2–3× larger than the true mask. Dice loss reduced this compared to CE-only (+7pp on both classes), but further improvement requires stronger spatial precision — achievable with attention mechanisms or a transformer-based decoder.
- Background class drag.** We compute 10-class mIoU (including Background = 32.38% IoU), while the paper reports 9-class mIoU. Background is a catch-all "unlabeled" class with noisy annotations, inherently limiting its achievable IoU. Our 9-class mIoU (59.37%) is a fairer comparison to the paper's metric.

5.2 Successful Improvements Over Baseline

Improvement	mIoU Gain	Key Benefit
50 epochs (vs 20)	+7.05pp	Model convergence
Dice loss added	(included above)	Reduced over-segmentation
LR Warmup (5 epochs)	+0.85pp	Eliminated catastrophic val loss spikes
75 epochs (vs 50)	+0.98pp	Better late-stage convergence
TTA (H+V flip)	+0.06pp	Orientation invariance at inference
Total vs 20-epoch baseline	+9.49pp	—

5.3 Future Work

To close the remaining ~8.6pp gap and potentially surpass GOA-U-Net:

- Stronger backbone (ResNet-50/101):** More feature channels improve boundary precision. Estimated +3–5pp mIoU. Requires re-running HHO search.
- SE/CBAM Attention in decoder:** Channel and spatial attention improve class-specific feature focus, directly addressing rare-class over-segmentation. Estimated +2–4pp.
- Transformer decoder (SegFormer):** Global self-attention enables cross-image context reasoning (e.g., distinguishing flooded road vs. water). Estimated +5–10pp, likely surpassing GOA-U-Net. Requires

additional libraries.

- **Longer training (100 epochs):** Model was still improving at epoch 74. Modest +0.5–1pp expected.

6. Conclusion

This project implements a complete end-to-end pipeline for HHO-optimized semantic segmentation of flood aerial imagery on the FloodNet dataset. Starting from a baseline of 47.18% mIoU, systematic improvements — extended training (75 epochs), combined CrossEntropy + Dice loss, linear learning rate warmup, and test-time augmentation — achieved a final **10-class mIoU of 56.67%** and **9-class mIoU of 59.37%**.

Our model surpasses 4 of 7 competing architectures in the reference paper and outperforms the best published model (GOA-U-Net) in **7 out of 9 classes by recall**, with exceptional gains on rare classes Vehicle (+55.6pp) and Pool (+39.3pp). The remaining mIoU gap is attributable to precision on rare classes rather than detection capability.

The HHO search converged to $\text{lr}=3.06\times10^{-4}$, $\text{batch}=8$, $\text{dropout}=0.189$, $\text{weight_decay}=5.70\times10^{-4}$ — consistent with the paper's findings ($\text{lr}\approx3\times10^{-4}$, $\text{batch}=4$). The larger optimal batch size (8 vs. 4) reflects the M4 Pro's unified memory advantage.

Appendix: Training Configuration

Parameter	Value
Architecture	ResNet-34 encoder + U-Net decoder
Total parameters	24,521,994
Input resolution	512 × 512 × 3
Output classes	10 (FloodNet)
Loss function	CrossEntropyLoss (weighted) + DiceLoss
Optimizer	Adam
Learning rate	3.060×10^{-4} (HHO-optimized)
LR schedule	Linear warmup (5ep) → CosineAnnealing ($\eta_{\min}=10^{-6}$)
Batch size	8 (HHO-optimized)
Dropout	0.189 (HHO-optimized)
Weight decay	5.699×10^{-4} (HHO-optimized)
Training epochs	75
Data augmentation	H-flip (50%), V-flip (50%), 90°/180°/270° rotation (50%)
Inference augmentation	TTA: average over original + H-flip + V-flip
Hardware	Apple M4 Pro (MPS backend)
Training time	≈ 3.1 hours
Random seed	42
Framework	PyTorch (Python 3.14)